

Model Free Prediction

岩延 yany@ucas.ac.cn

Model Free Reinforcement Learning

Introduction

- ▶ **Last Chapter**
 - Planning by dynamic programming
 - Solve a **known** MDP
- ▶ **Model Free prediction**
 - Estimate the value function of an **unknown** MDP
- ▶ **Model Free control**
 - Optimize the value function of an **unknown** MDP

Monte Carlo Methods

Introduction

- ▶ MC methods learn directly from episodes of experience
- ▶ MC is model-free: no knowledge of MDP transitions / rewards
- ▶ MC learns from complete episodes: no bootstrapping
- ▶ MC uses the simplest possible idea: value = mean return
- ▶ Can only apply MC to episodic MDPs
 - All episodes must terminate

Monte Carlo Methods

Monte Carlo Prediction

- ▶ 目标
 - 使用MC方法从经验中学习value function v_π
- ▶ Return is the total discounted reward:
 - $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1} R_T$
- ▶ Value function is the expected return:
 - $v_\pi(s) \doteq \mathbb{E}_\pi[G_t | S_t = s]$
- ▶ MC方法实质上就是利用回报的**Empirical Return**去代替**Expected Return**

Monte Carlo Methods

First-Visit Monte-Carlo

► To evaluate $v_\pi(s)$

- The **first** time-step t that state s is visited in an episode,
- Increment counter $N(s) \leftarrow N(s) + 1$
- Increment total return $S \leftarrow S + G_t$
- Value is estimated by mean return $V(s) = S/N(s)$

► 由大数定律可得，当 $N(s) \rightarrow \infty$

- $V(s) \rightarrow v_\pi(s)$

Monte Carlo Methods

Every-Visit Monte-Carlo

- ▶ To evaluate $v_\pi(s)$
 - Every time-step t that state s is visited in an episode,
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $S \leftarrow S + G_t$
 - Value is estimated by mean return $V(s) = S/N(s)$
- ▶ 由大数定律可得，当 $N(s) \rightarrow \infty$
 - $V(s) \rightarrow v_\pi(s)$

Monte Carlo Methods

Example

► 二十一点Black Jack


- 21点可以看成一个回合式的有限马尔科夫过程 (episodic finite MDP)
- 每次游戏都是一个回合 (episode)
- 赢、输、draw的奖励分别为1、-1、0
- 游戏过程中的任意动作奖励 (reward) 都为0
- No discount ($\gamma = 1$)
- 玩家的动作 (action) 只有要牌 (hits) 或者停止要牌 (sticks) 两种
- 如果玩家将A当成11点来算的话 (不能爆掉)，我们称它*usable*
- 总共有200个不同的状态

Monte Carlo Methods

Example


► 二十一点Black Jack

- 策略: 一直要牌, 直到点数和等于20或21时停止



Monte Carlo Methods

Backup Diagram of MC



Monte Carlo Methods

Incremental Implementation

▶ Incremental Mean

— 可以通过增量方式计算平均值

$$\begin{aligned} \bullet \quad Q_n &= \frac{R_1 + R_2 + \dots + R_{n-1}}{n-1} \\ \bullet \quad Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\ \bullet \quad &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \\ \bullet \quad &= \frac{1}{n} (R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i) \\ \bullet \quad &= \frac{1}{n} (R_n + (n-1)Q_n) \\ \bullet \quad &= \frac{1}{n} (R_n + nQ_n - Q_n) \\ \bullet \quad &= Q_n + \frac{1}{n} (R_n - Q_n) \end{aligned}$$

— 通用形式为

$$\bullet \quad \text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

Monte Carlo Methods

Incremental Monte-Carlo Updates

► To evaluate $v_\pi(s)$

- The first or every time-step t that state s is visited in an episode,
- Increment counter $N(s) \leftarrow N(s) + 1$
- Increment total return $S \leftarrow S + G_t$
- Value is estimated by mean return $V(s) = S/N(s)$

► 增量更新 $V(s)$

- For each state S_t with return G_t
 - $N(S_t) \leftarrow N(S_t) + 1$
 - $V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)}(G_t - V(S_t))$
- 对于non-stationary问题，可以使用固定的步长
 - $V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$

Temporal-Difference Learning

Introduction of Temporal-Difference Learning

- ▶ TD直接从episodes的经验中学习
- ▶ TD不需要完整的episodes
 - bootstrapping
- ▶ TD是model-free的
 - 不需要了解MDP transitions / rewards
- ▶ TD updates a guess towards a guess

Temporal-Difference Learning

TD Prediction

- ▶ 目标
 - 从经验中在线学习策略 π 的 v_π
- ▶ MC方法
 - 使用实际的 G_t 更新 $V(S_t)$
 - $V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$
- ▶ TD方法
 - 使用推测的return $R_{t+1} + \gamma V(S_{t+1})$ 更新 $V(S_t)$
 - $V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$
 - $R_{t+1} + \gamma V(S_{t+1})$ 为TD target
 - 称为TD(0)或one-step TD
- ▶ MC更新的Target是 G_t ， 而TD更新的Target是 $R_{t+1} + \gamma V(S_{t+1})$

Temporal-Difference Learning

TD Prediction

► 与DP类似， TD(0)是bootstrapping方法


- $v_\pi(s) \doteq E_\pi[G_t | S_t = s]$
- $\doteq E_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s]$
- $\doteq E_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s]$
- $\doteq \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$
- $\doteq \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$

► $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ 为 TD Error

- 存在多种形式

- MC error可以写成TD error之和

- $G_t - V(S_t) = R_{t+1} + \gamma G_{t+1} - V(S_t) + \gamma V(S_{t+1}) - \gamma V(S_{t+1})$
- $= \delta_t + \gamma(G_{t+1} - V(S_{t+1}))$
- $= \delta_t + \gamma\delta_{t+1} + \gamma^2(G_{t+2} - V(S_{t+2}))$
- $= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t}(G_T - V(S_T))$
- $= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t}(0)$
- $= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k$



Temporal-Difference Learning

Driving Home Example

状态	过去的时间	预测的时间	预测的总时间
周五六点，离开办公室	0	30	30
到车上，下雨	5	35	40
出主路	20	15	35
辅路，卡车后	30	10	40
进入居住街道	40	3	43
到家	43	0	43

Temporal-Difference Learning

Driving Home Example




Figure 6.1: Changes recommended in the driving home example by Monte Carlo methods (left) and TD methods (right).

Temporal-Difference Learning

TD Prediction

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal

Temporal-Difference Learning

MC vs TD

- ▶ **TD can learn **before** knowing the final outcome**
 - TD can learn **online** after every step
 - MC must wait until end of episode before return is known
- ▶ **TD can learn **without** the final outcome**
 - TD can learn from incomplete sequences
 - MC can only learn from complete sequences
 - TD works in continuing (non-terminating) environments
 - MC only works for episodic (terminating) environments

Temporal-Difference Learning

MC vs TD

- ▶ $v_\pi(S_t)$ 的 Unbiased 推测
 - $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1} R_T$
 - $R_{t+1} + \gamma v_\pi(S_{t+1})$
- ▶ $v_\pi(S_t)$ 的 Biased 推测
 - TD target $R_{t+1} + \gamma V(S_{t+1})$
- ▶ TD target 的 variance 比 return 低
 - Return depends on many random actions, transitions, rewards
 - TD target depends on one random action, transition, reward

Temporal-Difference Learning


MC vs TD

- ▶ **MC has high variance, zero bias**
 - Good convergence properties (even with function approximation)
 - Not very sensitive to initial value
 - Very simple to understand and use
- ▶ **TD has low variance, some bias**
 - Usually more efficient than MC
 - TD(0) converges to $v(s)$ (but not always with function approximation)
 - More sensitive to initial value

Temporal-Difference Learning

Example

► Random Walk



Temporal-Difference Learning

Batch MC and TD

- ▶ 随着经验趋于无穷MC和TD可保证收敛
 - $V(s) \rightarrow v_\pi(s)$
- ▶ 如果经验有限?
 - K 个episode
 - $s_1^1, a_1^1, r_1^1, \dots, s_{T_1}^1$
 - $s_1^k, a_1^k, r_1^k, \dots, s_{T_k}^k$
- ▶ Batch Update
 - 重复从这 K 个episode中进行采样
 - 对于某一次采样得到的episode $k \in [1, K]$ 应用MC或者TD(0)方法

Temporal-Difference Learning

Certainty-equivalence estimate

▶ Batch MC和TD的收敛性

- MC收敛于最小均方误差(minimum mean-squared error)解
 - $\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$
- TD收敛于极大似然Markov模型对应的解
 - MDP $\langle S, A, P, R, \gamma \rangle$ 为产生这些数据概率最大的模型
 - $p_{ss'}^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k, s_{t+1}^k = s, a, s')$
 - $r_s^a = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_t^k, a_t^k = s, a) r_t^k$

Temporal-Difference Learning

Example

- ▶ MC $V(A)=0$
- ▶ TD $V(A)=0.75$

Two states A, B ; no discounting; 8 episodes of experience

A, 0, B, 0

B, 1

B, 1


B, 1

B, 1

B, 1

B, 1

B, 0



What is $V(A), V(B)$?

Temporal-Difference Learning

Batch MC and TD


- ▶ **TD exploits Markov property**
 - Usually more efficient in Markov environments
- ▶ **MC does not exploit Markov property**
 - Usually more effective in non-Markov environments

Temporal-Difference Learning

Unified View

► MC backup

$$- V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$




Temporal-Difference Learning

Unified View

► TD backup

$$- V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$




Temporal-Difference Learning

Unified View

► DP backup

$$- V(s) = \mathbb{E}_\pi[R_{t+1} + \gamma V(S_{t+1})]$$



Temporal-Difference Learning

Unified View

► Bootstrapping and Sampling

- Bootstrapping: update involves an estimate
 - MC does not bootstrap
 - TD bootstraps
 - DP bootstraps
- Sampling: update samples an expectation
 - MC samples
 - TD samples
 - DP does not sample (it's full-width)

n-step Bootstrapping

n-step Bootstrapping

► MC

- wait until end of episode
 - $V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$
 - MC error $G_t - V(S_t)$

► TD(0)

- Wait until next time step
 - $V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$
 - TD error $R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$
 - TD target $R_{t+1} + \gamma V(S_{t+1})$


► n-step Bootstrapping

- 介于两者之间
- 称为n-step TD
- 利用eligibility traces

n-step Bootstrapping

n-step Prediction

- ▶ Let TD target look n steps into the future



n-step Bootstrapping

n-step Prediction

▶ 考虑 S_t 的State-value估计值

- 基于序列值 $S_t, R_{t+1}, S_{t+1}, R_{t+2}, \dots, R_T, S_T$
- MC: complete return
 - $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1} R_T$
- TD: one-step return
 - $G_{t:t+1} = R_{t+1} + \gamma V_t(S_{t+1})$
- two-step return
 - $G_{t:t+2} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V_{t+1}(S_{t+2})$

▶ n-step return

- Truncated
- Use n rewards and bootstrap
 - $G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$

▶ n-step TD learning

- $V_{t+n}(S_t) \leftarrow V_{t+n-1}(S_t) + \alpha(G_{t:t+n} - V_{t+n-1}(S_t))$

n-step Bootstrapping

n-step Prediction

n-step TD for estimating $V \approx v_\pi$

Input: a policy π

Algorithm parameters: step size $\alpha \in (0, 1]$, a positive integer n

Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}$

All store and access operations (for S_t and R_t) can take their index mod $n + 1$.

Loop for each episode:

Initialize and store $S_0 \neq \text{terminal}$

$$T \leftarrow \infty$$

Loop for $t = 0, 1, 2, \dots$:

If $t < T$, then:

Take an action according to $\pi(\cdot | S_t)$

Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

If S_{t+1} is terminal, then $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

If $\tau \geq 0$:

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$

$$V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$$

Until $\tau = T - 1$

n-step Bootstrapping

n-step Prediction

► Convergence


- The n-step return has **Error reduction property**
- 在最差情况下，n-step return相比 $V_{t+n-1}(S_{t+n})$ 是 $v_\pi(s)$ 的更好估计值
 - $\max_s |E_\pi[G_{t:t+n}|S_t = s] - v_\pi(s)| \leq \gamma^n \max_s |V_{t+n-1}(S_{t+n}) - v_\pi(s)|$
- 在合适的条件下可收敛至正确的value

n-step Bootstrapping

n-step Prediction

► Random Walk Example


- Rewards only on exit (-1 on left exit, 1 on right exit)
- n-step return: propagate reward up to n latest states



n-step Bootstrapping

n-step Prediction

Average
RMS error
over 19 states
and first 10
episodes



Eligibility traces

Eligibility Traces

- ▶ **Eligibility traces are one of the basic mechanisms of RL**
- ▶ **Eligibility traces unify and generalize TD and MC methods**
- ▶ **Advantage of eligibility traces over n-step methods**
 - Learning occurs continually and uniformly in time rather than being delayed and then catching up at the end of the episode.
 - Learning can occur and affect behavior immediately after a state is encountered rather than being delayed n steps.

Eligibility traces


Averaging n-Step Returns

- ▶ **n-step return**

- $G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$

- ▶ **We can average n-step returns over different n**

- e.g. average the 2-step and 4-step returns
 - $\frac{1}{2} G_{t:t+2} + \frac{1}{2} G_{t:t+4}$
 - Combines information from two different time-steps
 - Can we efficiently combine information from all time-steps?




Eligibility traces

λ -return

- The λ -return G_t^λ combines all n-step returns $G_{t:t+n}$

- Each has weight λ^{n-1}
- Normalized by $1 - \lambda$
- $G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$



- $(1 - \lambda) + (1 - \lambda)\lambda + (1 - \lambda)\lambda^2 + \dots + (1 - \lambda)\lambda^n = 1 + \lambda^n \approx 1$

Eligibility traces

λ -return

► TD(λ) Weighting Function

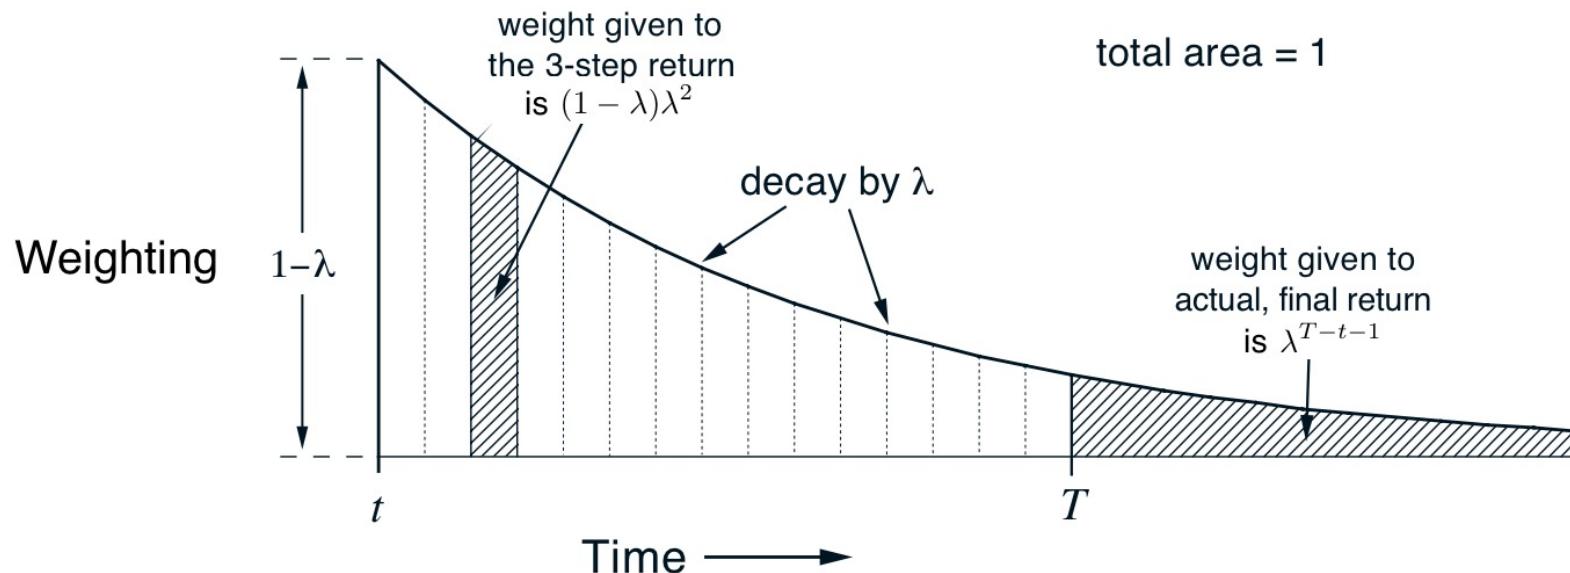


Figure 12.2: Weighting given in the λ -return to each of the n -step returns.

$$- G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_{t:t+n} + \lambda^{T-t-1} G_t$$

Eligibility traces

Forward-view TD(λ)

► Forward View TD(λ)

- $V(S_t) \leftarrow V(S_t) + \alpha(G_t^\lambda - V(S_t))$
- Update value function towards the λ -return
- Forward-view looks into the future to compute G_t^λ
- Like MC, can only be computed from complete episodes





Figure 12.4: The forward view. We decide how to update each state by looking forward to future rewards and states.

Eligibility traces

Forward-view TD(λ) on Large Random Walk



Eligibility traces


Backward-view TD(λ)

▶ Backward view

- Forward view provides theory
- Backward view provides mechanism
- Update online, every step, from incomplete sequences

▶ Eligibility Traces

- Credit assignment problem: did bell or light cause shock?
 - Frequency heuristic: assign credit to most frequent states
 - Recency heuristic: assign credit to most recent states
- Eligibility traces combine both heuristics
 - $E_0(s) = 0$
 - $E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$




Eligibility traces

Backward-view TD(λ)

► Backward-view TD(λ)

- Keep an eligibility trace for every state s
- Update value $V(s)$ for every state s
- In proportion to TD-error δ_t and eligibility trace $E_t(s)$
 - $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$
 - $V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$
- Compare with TD(0)
 - $V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$



Eligibility traces

Forwards and Backwards TD(λ)

- ▶ Consider an episode where s is visited once at time-step k
- ▶ TD(λ) eligibility trace discounts time since visit

- $E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s) = \begin{cases} 0 & \text{if } t < k \\ (\gamma\lambda)^{t-k} & \text{if } t \geq k \end{cases}$

- ▶ Backwards TD(λ) updates accumulate error online
 - $\sum_{t=1}^T \alpha \delta_t E_t(s) = \alpha \sum_{t=k}^T (\gamma\lambda)^{t-k} \delta_t = \alpha(G_k^\lambda - V(S_k))$
- ▶ By end of episode it accumulates total error for λ –return
- ▶ When $\lambda = 1$, sum of TD errors telescopes into MC error
 - $G_t - V(S_t) = R_{t+1} + \gamma G_{t+1} - V(S_t) + \gamma V(S_{t+1}) - \gamma V(S_{t+1})$
 - $= \delta_t + \gamma(G_{t+1} - V(S_{t+1}))$
 - $= \delta_t + \gamma\delta_{t+1} + \gamma^2(G_{t+2} - V(S_{t+2}))$
 - $= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \dots + \gamma^{T-t}(G_T - V(S_T))$
 - $= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \dots + \gamma^{T-t}(0)$
 - $= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k$

Eligibility traces

Offline Equivalence of Forward and Backward TD(λ)

► Offline updates

- Updates are accumulated within episode
- but applied in batch at the end of episode
- $\sum_{t=1}^T \alpha \left(G_t^\lambda - V(S_t) \right) \mathbf{1}(S_t = s) = \sum_{t=1}^T \alpha \delta_t E_t(s)$

► Online updates

- TD(λ) updates are applied online at each step within episode
- Forward and backward-view TD(λ) are slightly different

Eligibility traces

TD(λ) and TD(0)

- ▶ When $\lambda = 0$, only current state is updated
 - $E_t(s) = \mathbf{1}(S_t = s)$
 - $V(S) \leftarrow V(S) + \alpha \delta_t E_t(s)$
- ▶ This is exactly equivalent to TD(0) update
 - $V(S) \leftarrow V(S) + \alpha \delta_t = V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$

Eligibility traces

TD(1) and MC

- ▶ **TD(1) when $\lambda = 1$**
 - The credit given to earlier states falls only by γ per step
 - credit is deferred until end of episode
 - Achieve MC behavior
- ▶ **TD(1) can be applied to discounted continuing tasks**
- ▶ **TD(1) can be performed incrementally and online**
- ▶ **Over the course of an episode, total update for TD(1) is the same as total update for MC**

Eligibility traces

TD(1) and MC

- ▶ Consider an episode where s is visited once at time-step k
- ▶ TD(1) eligibility trace discounts time since visit
 - $E_t(s) = \gamma E_{t-1}(s) + \mathbf{1}(S_t = s) = \begin{cases} 0 & \text{if } t < k \\ \gamma^{t-k} & \text{if } t \geq k \end{cases}$
- ▶ TD(1) updates accumulate error online
 - $\sum_{t=1}^T \alpha \delta_t E_t(s) = \alpha \sum_{t=k}^T \gamma^{t-k} \delta_t = \alpha(G_k - V(S_k))$
- ▶ By end of episode it accumulates total error
 - $\sum_{t=k}^T \gamma^{t-k} \delta_t$

Eligibility traces

TD(1) and MC


- ▶ TD(1) is roughly equivalent to every-visit Monte-Carlo
- ▶ Error is accumulated online, step-by-step
- ▶ If value function is only updated offline at end of episode
 - Then total update is exactly the same as MC

Eligibility traces

Forward and Backward TD(λ)

Offline updates	$\lambda = 0$	$\lambda \in (0, 1)$	$\lambda = 1$
Backward view	TD(0) 	TD(λ) 	TD(1)
Forward view	TD(0)	Forward TD(λ)	MC
Online updates	$\lambda = 0$	$\lambda \in (0, 1)$	$\lambda = 1$
Backward view	TD(0) 	TD(λ) *	TD(1) *
Forward view	TD(0) 	Forward TD(λ) 	MC
Exact Online	TD(0)	Exact Online TD(λ)	Exact Online TD(1)

Unified View of Reinforcement Learning



Model Free Prediction

Suggested reading

- ▶ **Schultz, W., Dayan, P., and Montague, R., "A Neural Substrate of Prediction and Reward", 1999**
- ▶ **Redish, A. D., "Addiction as a Computational Process Gone Awry", 2004**