

Reinforcement Learning as Probabilistic Inference

岩延 yany@ucas.ac.cn

Variational Inference

Variational Inference



► Probabilistic models

- $p(x)$
- $p(y|x)$


► Latent variable models

- $p(x) = \sum_z p(x|z)p(z)$
- $= \int p(x|z)p(z)dz$
- $p(y|x) = \sum_z p(y|x, z)p(z)$

- $p(z)$
 - Easy distribution (e.g., Gaussian)
- $p(x|z)$
 - Easy distribution (e.g., Conditional Gaussian)



“easy” distribution
(e.g., Gaussian)



Variational Inference

Latent variable models



► Latent variable models

- Bayes rules


- $p(z|x) = \frac{p(x,z)}{p(x)}$

- Where

- $p(z|x)$ is posterior
- $p(z)$ is prior







“easy” distribution
(e.g., Gaussian)



Variational Inference

Variational Inference



Variational Inference

Variational Inference

- ▶ **The model**

- $p_\theta(x)$

- ▶ **The data**

- $D = \{x_1, x_2, \dots, x_N\}$

- ▶ **Maximum likelihood**

- $\theta \leftarrow \operatorname{argmax}_\theta \frac{1}{N} \sum_i \log p_\theta(x_i)$

- ▶ **Latent variable models**

- $p(x) = \int p(x|z)p(z)dz$

- ▶ **Maximum likelihood**

- $\theta \leftarrow \operatorname{argmax}_\theta \frac{1}{N} \sum_i \log(\int p_\theta(x_i|z)p(z)dz)$
 - Intractable

Variational Inference

Variational Inference

► Estimating the log-likelihood

- Guess the most likely z given x_i
- There are many possible values of z , so use the distribution $p(z|x_i)$
- How to calculate $p(z|x_i)$
 - Approximate with $q_i(z) = N(\mu_i, \sigma_i)$
 - Can be used to bound $\log p(x_i)$ for maximum likelihood

► We can bound $\log p(x_i)$

- $p(x_i) = \int p(x_i|z)p(z)dz$
- $\log p(x_i) = \log \int p(x_i|z)p(z)dz$
- $= \log \int p(x_i|z)p(z) \frac{q_i(z)}{q_i(z)} dz$
- $= \log \int q_i(z) \frac{p(x_i|z)p(z)}{q_i(z)} dz$
- $= \log \mathbb{E}_{z \sim q_i(z)} \left[\frac{p(x_i|z)p(z)}{q_i(z)} \right]$

Variational Inference

Variational Inference

- ▶ We can bound $\log p(x_i)$


- $\log p(x_i) = \log E_{z \sim q_i(z)} \left[\frac{p(x_i|z)p(z)}{q_i(z)} \right]$
- Use the Jensen's inequality
 - $\log E[y] \geq E[\log y]$
- $\log p(x_i) = \log E_{z \sim q_i(z)} \left[\frac{p(x_i|z)p(z)}{q_i(z)} \right]$
- $\geq E_{z \sim q_i(z)} \left[\log \frac{p(x_i|z)p(z)}{q_i(z)} \right]$
- $= E_{z \sim q_i(z)} [\log p(x_i|z) + \log p(z)] - E_{z \sim q_i(z)} [\log q_i(z)]$

Variational Inference

Variational Inference

► We can bound $\log p(x_i)$

- Remember the definition of Entropy
 - $H(p) = -\text{E}_{x \sim p(x)}[\log p(x)] = -\int p(x) \log p(x) dx$
 - How random is the random variable?
 - How large is the log probability in expectation under itself
- Then we have
 - $\log p(x_i) \geq \text{E}_{z \sim q_i(z)}[\log p(x_i|z) + \log p(z)] + H(q_i)$
 - ELBO (Evidence Lower Bound Objective)
 - $L_i(p, q_i) = \text{E}_{z \sim q_i(z)}[\log p(x_i|z) + \log p(z)] + H(q_i)$



Variational Inference

Variational Inference

► KL-Divergence Revisit

- $D_{KL}(q||p) = \text{E}_{x \sim q(x)}[\log \frac{q(x)}{p(x)}]$
- $= \text{E}_{x \sim q(x)}[\log q(x)] - \text{E}_{x \sim q(x)}[\log p(x)]$
- $= -\text{E}_{x \sim q(x)}[\log p(x)] - H(q)$
- How different are two distributions?
- How small is the expected log probability of one distribution under another, minus entropy?

Inverse Reinforcement Learning

Variational Approximation


► What makes a good $q_i(z)$?

- $q_i(z)$ should approximate $p(z|x_i)$
- in terms of the KL-divergence
 - $D_{KL}(q_i(z)||p(z|x_i))$
- $$D_{KL}(q_i(z)||p(z|x_i)) = \mathbb{E}_{z \sim q_i(z)} [\log \frac{q_i(z)}{p(z|x_i)}]$$
- $$= \mathbb{E}_{z \sim q_i(z)} [\log \frac{q_i(z)p(x_i)}{p(x_i|z)}] = \mathbb{E}_{z \sim q_i(z)} [\log \frac{q_i(z)p(x_i)}{p(x_i|z)p(z)}]$$
- $$= -\mathbb{E}_{z \sim q_i(z)} [\log p(x_i|z) + \log p(z)] + \mathbb{E}_{z \sim q_i(z)} [\log q_i(z)] + \mathbb{E}_{z \sim q_i(z)} [\log p(x_i)]$$
- $$= -\mathbb{E}_{z \sim q_i(z)} [\log p(x_i|z) + \log p(z)] - H(q_i) + \mathbb{E}_{z \sim q_i(z)} [\log p(x_i)]$$
 - Noticed that $\log p(x_i)$ independent of q_i
- $$= -L_i(p, q_i) + \log p(x_i)$$
- Then we have
 - $\log p(x_i) = D_{KL}(q_i(z)||p(z|x_i)) + L_i(p, q_i)$
 - $\log p(x_i) \geq L_i(p, q_i)$

Variational Inference

Variational Inference

- ▶ Maximizing $L_i(p, q_i)$ w.r.t. q_i minimizes KL-divergence
 - $\log p(x_i) = D_{KL}(q_i(z)||p(z|x_i)) + L_i(p, q_i)$



- Then our goal is to maximize ELBO $L_i(p, q_i)$
 - $L_i(p, q_i) = \mathbb{E}_{z \sim q_i(z)}[\log p(x_i|z) + \log p(z)] + H(q_i)$

Variational Inference

Variational Inference

► Maximum likelihood

- $L_i(p, q_i) = \mathbb{E}_{z \sim q_i(z)} [\log p_\theta(x_i|z) + \log p(z)] + H(q_i)$
- $\theta \leftarrow \operatorname{argmax}_\theta \frac{1}{N} \sum_i \log p_\theta(x_i)$ Change to
 - $\theta \leftarrow \operatorname{argmax}_\theta \frac{1}{N} \sum_i L_i(p, q_i)$
- For each x_i (or mini-batch)
 - Calculate $\nabla_\theta L_i(p, q_i)$
 - Sample $z \sim q_i(z)$
 - $\nabla_\theta L_i(p, q_i) \approx \nabla_\theta \log p_\theta(x_i|z)$
 - $\theta \leftarrow \theta + \alpha \nabla_\theta L_i(p, q_i)$
 - Update q_i to maximize $L_i(p, q_i)$
 - Let $q_i(z) = N(\mu_i, \sigma_i)$
 - Use gradient $\nabla_{\mu_i} L_i(p, q_i)$ and $\nabla_{\sigma_i} L_i(p, q_i)$ ascent on μ_i, σ_i

Variational Inference



Amortized variational inference

► What is the problem?

- Too many parameters
 - $|\theta| + (|\mu_i| + |\sigma_i|) \times N$

► How to solve?

- $q_i(z)$ should approximate $p(z|x_i)$
- Learn a network $q_i(z) = q(z|x_i) \approx p_\theta(x_i|z)$



- $\log p(x_i) \geq L_i(p_\theta(x_i|z), q_\phi(z|x_i))$
- $= \mathbb{E}_{z \sim q_\phi(z|x_i)} [\log p_\theta(x_i|z) + \log p(z)] + H(q_\phi(z|x_i))$
- For each x_i (or mini-batch)
 - Calculate $\nabla_\theta L_i(p_\theta(x_i|z), q_\phi(z|x_i))$
 - Sample $z \sim q_\phi(z|x_i)$
 - $\nabla_\theta L_i(p, q_i) \approx \nabla_\theta \log p_\theta(x_i|z)$
 - $\theta \leftarrow \theta + \alpha \nabla_\theta L_i$
 - $\phi \leftarrow \phi + \alpha \nabla_\phi L_i$
- How to calculate $\nabla_\phi L_i$?

Variational Inference

Amortized variational inference

► How to calculate $\nabla_{\phi} L_i$?

- $L_i = \mathbb{E}_{z \sim q_{\phi}(z|x_i)} [\log p_{\theta}(x_i|z) + \log p(z)] + H(q_{\phi}(z|x_i))$
- $q_{\phi}(z|x) = N(\mu_{\phi}(x), \sigma_{\phi}(x))$
 - $H(q_{\phi}(z|x_i))$ can be calculated easily
- Let $\log p_{\theta}(x_i|z) + \log p(z) = r(x_i, z)$
- Then we have
 - $J(\phi) = L_i = \mathbb{E}_{z \sim q_{\phi}(z|x_i)} [r(x_i, z)]$
 - Just use the policy gradient!
 - $J(\phi) \approx \frac{1}{M} \sum_j \nabla_{\phi} \log q_{\phi}(z_j|x_i) r(x_i, z_j)$
- What's the problem?
 - High variance

Variational Inference

Reparameterization

► How to calculate $J(\phi) = \mathbb{E}_{z \sim q_\phi(z|x_i)}[r(x_i, z)]$


- $q_\phi(z|x) = N(\mu_\phi(x), \sigma_\phi(x))$
- We can rewrite $z = \mu_\phi(x) + \epsilon\sigma_\phi(x)$ with noise $\epsilon \sim N(0,1)$
 - ϵ is independent to ϕ
- Then we can rewrite
 - $J(\phi) = \mathbb{E}_{z \sim q_\phi(z|x_i)}[r(x_i, z)]$
 - $= \mathbb{E}_{\epsilon \sim N(0,1)}[r(x_i, \mu_\phi(x_i) + \epsilon\sigma_\phi(x_i))]$
- Then calculate the $\nabla_\phi J(\phi)$
 - Sample $\epsilon_1, \epsilon_2, \dots, \epsilon_M$ from $N(0,1)$
 - $\nabla_\phi J(\phi) \approx \frac{1}{M} \sum_j \nabla_\phi r(x_i, \mu_\phi(x_i) + \epsilon_j\sigma_\phi(x_i))$

Variational Inference

Reparameterization

► Think in another way

- $L_i = \mathbb{E}_{z \sim q_\phi(z|x_i)} [\log p_\theta(x_i|z) + \log p(z)] + H(q_\phi(z|x_i))$
- $= \mathbb{E}_{z \sim q_\phi(z|x_i)} [\log p_\theta(x_i|z)] + \mathbb{E}_{z \sim q_\phi(z|x_i)} [\log p(z)] + H(q_\phi(z|x_i))$
- $= \mathbb{E}_{z \sim q_\phi(z|x_i)} [\log p_\theta(x_i|z)] - D_{KL}(q_\phi(z|x_i)||p(z))$
- $= \mathbb{E}_{\epsilon \sim N(0,1)} [\log p_\theta(x_i|\mu_\phi(x_i) + \epsilon\sigma_\phi(x_i))] - D_{KL}(q_\phi(z|x_i)||p(z))$
- $\approx \log p_\theta(x_i|\mu_\phi(x_i) + \epsilon\sigma_\phi(x_i)) - D_{KL}(q_\phi(z|x_i)||p(z))$



Variational Inference

Reparameterization

► Policy gradient

- $J(\phi) \approx \frac{1}{M} \sum_j \nabla_\phi \log q_\phi(z_j | x_i) r(x_i, z_j)$
- Can handle both discrete and continuous latent variables
- High variance, requires multiple samples & small learning rates




► Reparameterization trick

- $\nabla_\phi J(\phi) \approx \frac{1}{M} \sum_j \nabla_\phi r(x_i, \mu_\phi(x_i) + \epsilon_j \sigma_\phi(x_i))$
- Only continuous latent variables
- Very simple to implement
- Low variance

Variational Inference

The variational autoencoder

► It's a generative model



- $\max_{\theta, \phi} \frac{1}{N} \sum_i \log p_\theta(x_i | \mu_\phi(x_i) + \epsilon\sigma_\phi(x_i)) - D_{KL}(q_\phi(z|x_i) || p(z))$
- $p(x) = \int p(x|z)p(z)dz$
 - Sampling $z \sim p(z)$ $x \sim p(x|z)$

Probabilistic Graphical Model



Optimal Control

- ▶ Does reinforcement learning and optimal control provide a reasonable model of human behavior?

- $a_1, a_2, \dots, a_T = \underset{a_1, a_2, \dots, a_T}{\operatorname{argmax}} \sum_{t=1}^T r(s_t, a_t)$
- $s_{t+1} = p(s_t, a_t)$
- $\pi = \underset{\pi}{\operatorname{argmax}} \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t), a_t \sim \pi(s_t | a_t)} [r(s_t, a_t)]$
- $a_t \sim \pi(s_t | a_t)$

- ▶ What if the data is not optimal?

- some mistakes matter more than others
- behavior is stochastic
- but good behavior is still the most likely




Probabilistic Graphical Model

Probabilistic Graphical Model

► Probabilistic Graphical Model

- $p(\tau) = p(s_{1:T}, a_{1:T}) = ?$? No assumption of optimal behavior
- $p(O_t|s_t, a_t) \propto \exp(r(s_t, a_t))$
- $p(\tau|O_{1:T}) = \frac{p(\tau, O_{1:T})}{p(O_{1:T})} \propto p(\tau, O_{1:T})$
- $= p(s_1) \prod_t p(O_t|s_t, a_t) p(s_{t+1}|s_t, a_t)$
- $= p(s_1) \prod_t \exp(r(s_t, a_t)) p(s_{t+1}|s_t, a_t)$
- $= p(\tau) \prod_t \exp(r(s_t, a_t))$
- $= p(\tau) \exp(\sum_t r(s_t, a_t))$
- Optimal policy
 - $p(a_t|s_t, O_{t:T} = 1)$




Probabilistic Graphical Model

Probabilistic Graphical Model

► Why Probabilistic Graphical Model?

- Can model suboptimal behavior (important for inverse RL)
- Can apply inference algorithms to solve control and planning problems
- Provides an explanation for why stochastic behavior might be preferred (useful for exploration and transfer learning)




Probabilistic Graphical Model

Inference = Control

► How to do inference?

- Compute backward message $\beta_t(s_t, a_t) = p(O_{t:T} | s_t, a_t)$
- Compute policy $p(a_t | s_t, O_{1:T})$
- Compute forward messages $\alpha_t(s_t) = p(s_t | O_{1:t-1})$



Probabilistic Graphical Model

Inference = Control

► Backward messages

- $\beta_t(s_t, a_t) = p(O_{t:T} | s_t, a_t)$
- $= \int p(O_{t:T}, s_{t+1} | s_t, a_t) ds_{t+1}$
- $= \int p(O_{t+1:T} | s_{t+1}) p(s_{t+1} | s_t, a_t) p(O_t | s_t, a_t) ds_{t+1}$
- Let $\beta_t(s_t) = p(O_{t:T} | s_t)$, we have
 - $\beta_t(s_t, a_t) = \int \beta_{t+1}(s_{t+1}) p(s_{t+1} | s_t, a_t) p(O_t | s_t, a_t) ds_{t+1}$
 - $\beta_t(s_t) = p(O_{t:T} | s_t)$
 - $= \int p(O_{t:T} | s_t, a_t) p(a_t | s_t) da_t$
 - $= \int \beta_t(s_t, a_t) p(a_t | s_t) da_t$
 - Where $p(O_{t:T} | s_t, a_t) = \beta_t(s_t, a_t)$
- Then we have
 - $\beta_T(s_T, a_T) = p(O_T | s_t, a_t) \propto \exp(r(s_T, a_T))$
 - From $t = T - 1$ to 1
 - $\beta_t(s_t, a_t) = p(O_t | s_t, a_t) \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} [\beta_{t+1}(s_{t+1})]$
 - $\beta_t(s_t) = \mathbb{E}_{a_t \sim p(a_t | s_t)} [\beta_t(s_t, a_t)]$

Probabilistic Graphical Model

Backward Messages vs Value Iteration

► Backward messages

- From $t = T - 1$ to 1
 - $\beta_t(s_t, a_t) = p(O_t | s_t, a_t) \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} [\beta_{t+1}(s_{t+1})]$
 - $\beta_t(s_t) = \mathbb{E}_{a_t \sim p(a_t | s_t)} [\beta_t(s_t, a_t)]$
- Log space message
 - Let $V(s_t) = \log \beta_t(s_t)$, then $\beta_t(s_t) = \exp(V(s_t))$
 - Let $Q(s_t, a_t) = \log \beta_t(s_t, a_t)$, then $\beta_t(s_t, a_t) = \exp(Q(s_t, a_t))$
- Then we have
 - $V(s_t) = \log \int \exp(Q(s_t, a_t)) da_t$
 - $V(s_t) \rightarrow \max_{a_t} Q(s_t, a_t)$ as $Q(s_t, a_t)$ gets bigger
 - Means **soft max!**
 - $Q(s_t, a_t) = \log p(O_t | s_t, a_t) \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} [\beta_{t+1}(s_{t+1})]$
 - $= r(s_t, a_t) + \log \mathbb{E}[\beta_{t+1}(s_{t+1})]$
 - $= r(s_t, a_t) + \log \mathbb{E}[\exp(V(s_{t+1}))]$
 - For deterministic transition we have
 - $Q(s_t, a_t) = r(s_t, a_t) + V(s_{t+1})$

Probabilistic Graphical Model

Backward Messages vs Value Iteration

► Remember value iteration

- $V_{k+1}(s) = \max_a \mathbb{E}[R_{t+1} + \gamma V_k(S_{t+1}) | S_t = s, A_t = a]$
- $V(s) \leftarrow \max_a Q(s, a)$
- $Q(s, a) \leftarrow r(s, a) + \gamma \mathbb{E}[V(s')]$

► Backward messages


- $V(s_t) = \log \int \exp(Q(s_t, a_t)) da_t \rightarrow \max_{a_t} Q(s_t, a_t)$
- $Q(s_t, a_t) = r(s_t, a_t) + \log \mathbb{E}[\exp(V(s_{t+1}))]$
- $\log \mathbb{E}[\exp(V(s_{t+1}))]$ is an optimistic transition, not a good idea
- A better stochastic model
 - $Q(s_t, a_t) = r(s_t, a_t) + \mathbb{E}[V(s_{t+1})]$

Probabilistic Graphical Model

Backward Pass Summary

► Backward messages

- $\beta_t(s_t, a_t) = p(O_{t:T} | s_t, a_t)$
 - Probability that we can be optimal at step t through T given that we take action a_t in state s_t
- From $t = T - 1$ to 1
 - $\beta_t(s_t, a_t) = p(O_t | s_t, a_t) \mathbb{E}_{s_{t+1} \sim p(s_{t+1} | s_t, a_t)} [\beta_{t+1}(s_{t+1})]$
 - $\beta_t(s_t) = \mathbb{E}_{a_t \sim p(a_t | s_t)} [\beta_t(s_t, a_t)]$
- Log space message
 - Let $V(s_t) = \log \beta_t(s_t)$
 - Let $Q(s_t, a_t) = \log \beta_t(s_t, a_t)$
 - Log of β_t is Q-function-like



Probabilistic Graphical Model

The Action Prior

- ▶ We assumed uniform action prior

- $p(O_{t+1:T}|s_{t+1}) = \beta_{t+1}(s_{t+1})$
- $= \int p(O_{t+1:T}|s_{t+1}, a_{t+1})p(a_{t+1}|s_{t+1})da_{t+1}$

- ▶ What if action prior is not uniform?


- $V(s_t) = \log \int \exp(Q(s_t, a_t)) + \log p(a_t|s_t) da_t$
- $Q(s_t, a_t) = r(s_t, a_t) + \log E[\exp(V(s_{t+1}))]$
- Let
 - $\tilde{Q}(s_t, a_t) = r(s_t, a_t) + \log p(a_t|s_t) + \log E[\exp(V(s_{t+1}))]$
- Then we have
 - $V(s_t) = \log \int \exp(\tilde{Q}(s_t, a_t)) da_t$
- Can always fold the action prior into the reward! uniform action prior can be assumed without loss of generality

Probabilistic Graphical Model

Policy Computation

► Compute policy $p(a_t|s_t, O_{1:T})$

$$\begin{aligned} - \beta_t(s_t, a_t) &= p(O_{t:T}|s_t, a_t) & \beta_t(s_t) &= p(O_{t:T}|s_t) \\ - p(a_t|s_t, O_{1:T}) &= \pi(a_t|s_t) = p(a_t|s_t, O_{t:T}) \\ - &= \frac{p(a_t, s_t | O_{t:T})}{p(s_t | O_{t:T})} \\ - &= \frac{p(O_{t:T} | a_t, s_t) p(a_t, s_t) / p(O_{t:T})}{p(O_{t:T} | s_t) p(s_t) / p(O_{t:T})} = \frac{p(O_{t:T} | a_t, s_t) p(a_t, s_t)}{p(O_{t:T} | s_t) p(s_t)} \\ - &= \frac{\beta_t(s_t, a_t)}{\beta_t(s_t)} p(a_t | s_t) = \frac{\beta_t(s_t, a_t)}{\beta_t(s_t)} \end{aligned}$$



Probabilistic Graphical Model

Policy Computation

► Policy computation with value functions

- From $t = T - 1$ to 1
 - $Q(s_t, a_t) = r(s_t, a_t) + \log E[\exp(V(s_{t+1}))]$
 - $V(s_t) = \log \int \exp(Q(s_t, a_t)) da_t$
- $\pi(a_t | s_t) = \frac{\beta_t(s_t, a_t)}{\beta_t(s_t)}$
- $V(s_t) = \log \beta_t(s_t)$
- $Q(s_t, a_t) = \log \beta_t(s_t, a_t)$
- Then we have
 - $\pi(a_t | s_t) = \exp(Q(s_t, a_t) - V(s_t))$
 - $= \exp(A(s_t, a_t))$
 - Better action are more probable

Probabilistic Graphical Model

Forward Messages

► Forward messages

- $\alpha_1(s_1) = p(s_1)$
- $\alpha_t(s_t) = p(s_t | O_{1:t-1})$
- $= \int p(s_t | s_{t-1}, a_{t-1}) p(a_{t-1} | s_{t-1}, O_{t-1}) p(s_{t-1} | O_{1:t-2}) ds_{t-1} da_{t-1}$
- $= \int p(s_t | s_{t-1}, a_{t-1}) p(a_{t-1} | s_{t-1}, O_{t-1}) \alpha_{t-1}(s_{t-1}) ds_{t-1} da_{t-1}$
- Where
 - $p(a_{t-1} | s_{t-1}, O_{t-1}) = \frac{p(O_{t-1} | s_{t-1}, a_{t-1}) p(a_{t-1} | s_{t-1})}{p(O_{t-1} | s_{t-1})}$
 - Remember $p(O_{t-1} | s_{t-1}, a_{t-1}) \propto \exp(r(s_{t-1}, a_{t-1}))$
 - Then $p(a_{t-1} | s_{t-1}, O_{t-1}) \propto \exp(r(s_{t-1}, a_{t-1}))$

► We want $p(s_t | O_{1:T})$

- $p(s_t | O_{1:T}) = \frac{p(s_t, O_{1:T})}{p(O_{1:T})} = \frac{p(O_{1:T} | s_t) p(s_t, O_{1:t-1})}{p(O_{1:T})}$
- $\propto \beta_t(s_t) \alpha_t(s_t)$



Probabilistic Graphical Model

Forward/backward message intersection

states with high probability of reaching goal

$$p(s_t) \propto \beta_t(s_t)\alpha_t(s_t)$$

states with high probability of being reached from initial state (with high reward)



Probabilistic Graphical Model

Summary

- ▶ Probabilistic graphical model for optimal control
- ▶ Control = inference
- ▶ Very similar to dynamic programming, value iteration, etc.
(but “soft”)

Probabilistic Graphical Model

The Optimism Problem

► The optimism issue

- $Q(s_t, a_t) = r(s_t, a_t) + \log E[\exp(V(s_{t+1}))]$
- $\log E[\exp(V(s_{t+1}))]$ is an optimistic transition, not a good idea

► Inference problem

- $p(s_{1:T}, a_{1:T} | O_{1:T})$
- We get policy $p(a_t | s_t, O_{1:T})$ by marginalizing and conditioning
 - Means given that you obtained high reward, what was your action probability?
- we can also get $p(s_{t+1} | s_t, a_t, O_{1:T})$ by marginalizing and conditioning
 - $p(s_{t+1} | s_t, a_t, O_{1:T}) \neq p(s_{t+1} | s_t, a_t)$
 - Means given you obtained high reward, what was your transition probability?
- We want to get right action $p(a_t | s_t, O_{1:T})$

Probabilistic Graphical Model

The Optimism Problem

► The optimism problem

- given that you obtained high reward, what was your action probability, given that your transition probability didn't change?

► How can we solve that?



- Find another distribution $q(s_{1:T}, a_{1:T})$ that is close to $p(s_{1:T}, a_{1:T}|O_{1:T})$ but has dynamics $p(s_{t+1}|s_t, a_t)$
- Can be solved by variational inference
 - Let $x = O_{1:T}$
 - Let $z = s_{1:T}, a_{1:T}$
 - The problem is equal to
 - Find $q(z)$ to approximate $p(z|x)$

Probabilistic Graphical Model

Control via Variational Inference

► How to solve it?

- Let $q(z) = q(s_{1:T}, a_{1:T}) = q(\tau) = p(s_1) \prod_t p(s_{t+1}|s_t, a_t) \pi(a_t|s_t)$
- Let $p(z|x) = p(s_{1:T}, a_{1:T}|O_{1:T}) = p(\tau|O_{1:T})$
- Remember $p(\tau|O_{1:T}) = \frac{p(\tau, O_{1:T})}{p(O_{1:T})} \propto p(\tau, O_{1:T}) = p(x, z)$
 - $= p(s_1) \prod_t p(O_t|s_t, a_t) p(s_{t+1}|s_t, a_t)$
 - $= p(\tau) \prod_t \exp(r(s_t, a_t))$
 - $= p(\tau) \exp(\sum_t r(s_t, a_t))$



Probabilistic Graphical Model

Control via Variational Inference

► The variational lower bound

– Remember

- $\log p(x) \geq \mathbb{E}_{z \sim q(z)}[\log p(x, z) - \log q(z)]$

– Then

- $\log p(x) = \log p(O_{1:T})$
- $\geq \mathbb{E}_{(s_{1:T}, a_{1:T}) \sim \pi}[\log p(s_1) + \sum_{t=1}^T \log p(s_{t+1}|s_t, a_t) + \sum_{t=1}^T \log p(O_t|s_t, a_t)]$
- $\quad - \log p(s_1) - \sum_{t=1}^T \log p(s_{t+1}|s_t, a_t) - \sum_{t=1}^T \log \pi(a_t|s_t)]$
- $= \mathbb{E}_{(s_{1:T}, a_{1:T}) \sim \pi}[\sum_t r(s_t, a_t) - \log \pi(a_t|s_t)]$
- $= \sum_t \mathbb{E}_{(s_t, a_t) \sim \pi}[r(s_t, a_t) + H(\pi(a_t|s_t))]$

▪ **Maximize reward and action entropy**

▪ Temperature could be used

- $\sum_t \mathbb{E}_{(s_t, a_t) \sim \pi}[r(s_t, a_t) + \alpha H(\pi(a_t|s_t))]$

Probabilistic Graphical Model

Control via Variational Inference

► Optimizing the variational lower bound

- $\log p(x) = \log p(O_{1:T}) \geq \sum_t \mathbb{E}_{(s_t, a_t) \sim \pi} [r(s_t, a_t) + H(\pi(a_t | s_t))]$
-
- Solve for $\pi(a_T | s_T)$ first
 - $\pi(a_T | s_T) = \operatorname{argmax}_{s_T \sim q(s_T)} [\mathbb{E}_{a_T \sim \pi(a_T | s_T)} [r(s_T, a_T)] + H(\pi(a_T | s_T))]$
 - $= \operatorname{argmax}_{s_T \sim q(s_T)} [\mathbb{E}_{a_T \sim \pi(a_T | s_T)} [r(s_T, a_T) - \log \pi(a_T | s_T)]]$
 - Solved when $\pi(a_T | s_T) \propto \exp(r(s_T, a_T))$
 - $\pi(a_T | s_T) = \frac{\exp(r(s_T, a_T))}{\int \exp(r(s_T, a)) da} = \frac{\exp(r(s_T, a_T))}{\exp(V(s_T))} = \exp(Q(s_T, a_T) - V(s_T))$
 - Where $V(s_T) = \log \int \exp(Q(s_T, a_T)) da_T = \log \int \exp(r(s_T, a_T)) da_T$
 - Then we have
 - $\mathbb{E}_{s_T \sim q(s_T)} [\mathbb{E}_{a_T \sim \pi(a_T | s_T)} [r(s_T, a_T) - \log \pi(a_T | s_T)]]$
 - $= \mathbb{E}_{s_T \sim q(s_T)} [\mathbb{E}_{a_T \sim \pi(a_T | s_T)} [V(s_T)]]$

Probabilistic Graphical Model

Control via Variational Inference

► Solve recursively

- $\pi(a_t|s_t) = \operatorname{argmax} E_{s_t, a_t \sim \pi}[r(s_t, a_t) + H(\pi(a_t|s_t)) + E_{s_{t+1} \sim p(s_{t+1}|s_t, a_t)}[V(s_{t+1})]]$
- Now define $Q(s_t, a_t) = r(s_t, a_t) + E[V(s_{t+1})]$
 - Now it's regular Bellman backup
 - A better stochastic model
- $\pi(a_t|s_t) = \operatorname{argmax} E_{s_t \sim q(s_t)}[E_{a_t \sim \pi(a_t|s_t)}[Q(s_t, a_t)] + H(\pi(a_t|s_t))]$
- $= \operatorname{argmax} E_{s_t \sim q(s_t)}[E_{a_t \sim \pi(a_t|s_t)}[Q(s_t, a_t) - \log \pi(a_t|s_t)]]$
 - Solved when $\pi(a_t|s_t) \propto \exp(r(s_t, a_t))$
- $\pi(a_t|s_t) = \exp(Q(s_t, a_t) - V(s_t))$
- $V(s_t) = \log \int \exp(Q(s_t, a_t)) da_t = E_{a_t \sim \pi}[Q(s_t, a_t) - \log \pi(a_t|s_t)]$
- $Q(s_t, a_t) = r(s_t, a_t) + E[V(s_{t+1})]$

Probabilistic Graphical Model

Backward pass summary - variational

▶ Backward pass

- From $t = T - 1$ to 1
 - $Q(s_t, a_t) = r(s_t, a_t) + \mathbb{E}[V(s_{t+1})]$
 - $V(s_t) = \log \int \exp(Q(s_t, a_t)) da_t$

▶ Value iteration

- $Q(s, a) \leftarrow r(s, a) + \gamma \mathbb{E}[V(s')]$
- $V(s) \leftarrow \max_a Q(s, a)$

▶ Soft value iteration

- $Q(s, a) \leftarrow r(s, a) + \gamma \mathbb{E}[V(s')]$
- $V(s) \leftarrow \text{soft max}_a Q(s, a)$

▶ Discount and temperature

- Discounted SOC $Q(s_t, a_t) = r(s_t, a_t) + \gamma \mathbb{E}[V(s_{t+1})]$
- Explicit temperature $V(s_t) = \alpha \log \int \exp(\frac{1}{\alpha} Q(s_t, a_t)) da_t$
 - With temperature $\pi(a_t | s_t) = \exp(\frac{1}{\alpha} Q(s_t, a_t) - \frac{1}{\alpha} V(s_t))$
 - $= \exp(\frac{1}{\alpha} A(s_t, a_t))$

Probabilistic Graphical Model

Q-learning with soft optimality

► Q-learning

- $\phi \leftarrow \phi + \alpha \nabla_{\phi} Q_{\phi}(s, a)(r(s, a) + \gamma V(s') - Q_{\phi}(s, a))$
- Target: $V(s') = \max_{a'} Q_{\phi}(s', a')$

► Q-learning with soft optimality

- Soft Q-learning: $\phi \leftarrow \phi + \alpha \nabla_{\phi} Q_{\phi}(s, a)(r(s, a) + \gamma V(s') - Q_{\phi}(s, a))$
- Target: $V(s') = \text{soft max}_{a'} Q_{\phi}(s', a') = \log \int \exp(Q_{\phi}(s', a')) da'$
- $\pi(a|s) = \exp(Q_{\phi}(s, a) - V(s)) = \exp(A(s, a))$

► Soft DQN

- Take some action a_i and observe (s_i, a_i, s'_i, r_i) , add it to \mathcal{B}
- Sample mini-batch of transitions $\{(s_i, a_i, s'_i, r_i)\}$ from \mathcal{B} uniformly
- Compute $y_i = r(s_i, a_i) + \gamma \text{soft max}_{a'_i} Q_{\phi'}^{\pi}(s'_i, a'_i)$ using target network $Q_{\phi'}^{\pi}$
- $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_{\phi}^{\pi}}{d\phi}(s_i, a_i)(Q_{\phi}^{\pi}(s_i, a_i) - y_i)$
- Update ϕ' : copy ϕ every N steps

Probabilistic Graphical Model

Policy gradient with soft optimality

► Entropy regularized policy gradient

- When $\pi(a|s) = \exp(Q_\phi(s, a) - V(s))$
 - It optimize $\sum_t \mathbb{E}_{\pi(s_t, a_t)}[r(s_t, a_t)] + \mathbb{E}_{\pi(s_t)}[H(\pi(a_t|s_t))]$
 - $H(\pi(a_t|s_t))$ is the policy entropy
- $\pi(a|s) \propto \exp(Q_\phi(s, a))$ when π minimizes
 $D_{KL}(\pi(a|s) || \frac{1}{Z} \exp(Q_\phi(s, a)))$
 - $D_{KL}(\pi(a|s) || \frac{1}{Z} \exp(Q_\phi(s, a))) = \mathbb{E}_{\pi(a|s)}[Q_\phi(s, a)] + H(\pi)$
- Entropy regularized policy gradient combats premature entropy collapse
- Closely related to soft Q-learning

Probabilistic Graphical Model

Soft Actor-Critic

► Soft Policy Iteration

- Still optimize $\sum_t \mathbb{E}_{\pi(s_t, a_t)}[r(s_t, a_t)] + \mathbb{E}_{\pi(s_t)}[H(\pi(a_t | s_t))]$
- Evaluation
 - $Q(s_t, a_t) = r(s_t, a_t) + \gamma \mathbb{E}_{s_{t+1} \sim p}[V(s_{t+1})]$
 - $V(s_t) = \mathbb{E}_{a_t \sim \pi}[Q(s_t, a_t) - \alpha \log \pi(a_t | s_t)]$
- Improvement
 - $\pi_{\text{new}} = \operatorname{argmin}_{\pi' \in \Pi} D_{KL}(\pi'(\cdot | s) || \frac{\exp(Q_{\pi^{\text{old}}}(s_t, \cdot))}{Z_{\pi^{\text{old}}}(s_t)})$

Probabilistic Graphical Model

Soft Actor-Critic

► Soft Actor-Critic

- Use neural network $Q_\theta(s_t, a_t)$, $\pi_\phi(a_t | s_t)$, $V_\psi(s_t)$
- Train value function
 - $J_\psi(V) = \mathbb{E}_{s_t \sim D} [\frac{1}{2} (V_\psi(s_t) - \mathbb{E}_{a_t \sim \pi_\phi} [Q_\theta(s_t, a_t) - \alpha \log \pi_\phi(a_t | s_t)])^2]$
- Train Q function
 - $J_\theta(Q) = \mathbb{E}_{s_t, a_t \sim D} [\frac{1}{2} (Q_\theta(s_t, a_t) - \hat{Q}_\theta(s_t, a_t))^2]$
 - Where $\hat{Q}_\theta(s_t, a_t) = r(s_t, a_t) + \gamma \mathbb{E}_{s_{t+1} \sim p} [V_\psi(s_{t+1})]$
- Train Policy
 - $J_\phi(\pi) = \mathbb{E}_{s_t \sim D} [D_{KL}(\pi_\phi(\cdot | s_t) || \frac{\exp(Q_\theta(s_t, \cdot))}{Z_\theta(s_t)})]$
- Reparametrize trick
 - $a_t = f_\phi(\epsilon_t; s_t)$
 - $J_\phi(\pi) = \mathbb{E}_{s_t \sim D, \epsilon_t \sim N} [D_{KL}(\pi_\phi(\cdot | s_t) || \frac{\exp(Q_\theta(s_t, \cdot))}{Z_\theta(s_t)})]$

Probabilistic Graphical Model

Soft Actor-Critic

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$$

$$\bar{\psi} \leftarrow \tau\psi + (1 - \tau)\bar{\psi}$$

end for

end for

Probabilistic Graphical Model

Review

- ▶ **Reinforcement learning can be viewed as inference in a graphical model**
 - Value function is a backward message
 - Maximize reward and entropy (the bigger the rewards, the less entropy matters)
 - Variational inference to remove optimism
- ▶ **Soft Q-learning**
- ▶ **Entropy-regularized policy gradient**
- ▶ **Soft Actor-Critic**

Probabilistic Graphical Model

Benefits of soft optimality

- ▶ Improve exploration and prevent entropy collapse
- ▶ Easier to specialize (finetune) policies for more specific tasks
- ▶ Principled approach to break ties
- ▶ Better robustness (due to wider coverage of states)
- ▶ Can reduce to hard optimality as reward magnitude increases
- ▶ Good model for modeling human behavior (more on this later)

Inverse Reinforcement Learning

Inverse Reinforcement Learning

- ▶ **So far**
 - Manually design reward function to define a task
- ▶ **What if we want to learn the reward function from observing an expert, and then use reinforcement learning?**
 - Apply approximate optimality model from last week, but now learn the reward
- ▶ **Inverse reinforcement learning**
 - Infer reward function from roll-outs of expert policy

Inverse Reinforcement Learning

Inverse Reinforcement Learning

► Why should we learn the reward?

- Alternative: directly mimic the expert (behavior cloning)
 - simply “ape” the expert’s motions/actions
 - doesn’t necessarily capture the salient parts of the behavior
 - what if the expert has different capabilities?
- Can we reason about what the expert is trying to achieve instead?

Inverse Reinforcement Learning

Inverse Optimal Control / Inverse Reinforcement Learning:

► **Infer reward function from demonstrations**

- Given
 - State & action space
 - Samples from π^*
 - Dynamics model (sometimes)
- Goal
 - Recover reward function
 - Then use reward to get policy
- Challenges
 - Underdefined problem
 - Difficult to evaluate a learned reward
 - Demonstrations may not be precisely optimal

Inverse Reinforcement Learning

Inverse Reinforcement Learning

► Forward Reinforcement Learning

- Given
 - States $s \in S$ actions $a \in A$
 - (sometimes) Transitions $p(s'|s, a)$
 - Reward function $r(s, a)$
- Learn $\pi^*(a|s)$

► Inverse Reinforcement Learning

- Given
 - States $s \in S$ actions $a \in A$
 - (sometimes) transitions $p(s'|s, a)$
 - Samples $\{\tau_i\}$ from $\pi^*(\tau)$
- Learn $r_\psi(s, a)$
 - Could be linear $r_\psi(s, a) = \sum_i \psi_i f_i(s, a) = \psi^T f(s, a)$
 - Or neural net
- Then use it to learn $\pi^*(a|s)$

Inverse Reinforcement Learning

Feature matching IRL

- ▶ **Linear reward function**

- $r_\psi(s, a) = \sum_i \psi_i f_i(s, a) = \psi^T f(s, a)$

- ▶ **If the feature f is important, we could match their expectations**

- Let π^{r_ψ} be the optimal policy for r_ψ
 - Pick ψ such that $E_{\pi^{r_\psi}}[f(s, a)] = E_{\pi^*}[f(s, a)]$
 - $E_{\pi^{r_\psi}}[f(s, a)]$ is the state-action marginal under π^{r_ψ}
 - $E_{\pi^*}[f(s, a)]$ is unknown optimal policy approximate using expert samples
 - Using maximum margin principle
 - $\max_{\psi, m} m$
 - Such that
 - $\psi^T E_{\pi^*}[f(s, a)] \geq \max_{\pi \in \Pi} \psi^T E_\pi[f(s, a)] + m$

Inverse Reinforcement Learning

Feature matching IRL & maximum margin


► Remember the SVM

- $\max_{\psi, m} m$ Such that $\psi^T E_{\pi^*}[f(s, a)] \geq \max_{\pi \in \Pi} \psi^T E_\pi[f(s, a)] + m$
- $\min_{\psi} \frac{1}{2} \|\psi\|^2$ such that $\psi^T E_{\pi^*}[f(s, a)] \geq \max_{\pi \in \Pi} \psi^T E_\pi[f(s, a)] + D(\pi, \pi^*)$
- Where $D(\pi, \pi^*)$ is difference in feature expectations
- Problems
 - Maximizing the margin is a bit arbitrary
 - No clear model of expert suboptimality (can add slack variables...)
 - Messy constrained optimization problem – not great for deep learning

Inverse Reinforcement Learning

Probabilistic Graphical Model

- ▶ A probabilistic graphical model of decision making



- $p(s_{1:T}, a_{1:T}) = ??$ $p(\mathcal{O}_t | s_t, a_t) \propto \exp(r(s_t, a_t))$
- $p(\tau | \mathcal{O}_{1:T}) = \frac{p(\tau, \mathcal{O}_{1:T})}{p(\mathcal{O}_{1:T})} \propto p(\tau) \prod_t \exp(r(s_t, a_t)) = p(\tau) \exp(\sum_t r(s_t, a_t))$

Inverse Reinforcement Learning

Probabilistic Graphical Model

▶ Learning the optimality variable

- $p(O_t|s_t, a_t, \psi) \propto \exp(r_\psi(s_t, a_t))$
- $p(\tau|O_{1:T}, \psi) = \frac{p(\tau, O_{1:T})}{p(O_{1:T})} \propto p(\tau) \prod_t \exp(r_\psi(s_t, a_t)) \sim \frac{1}{Z} p(\tau) \exp(r_\psi(\tau))$
- Given
 - Samples $\{\tau_i\}$ from $\pi^*(\tau)$
 - We can perform maximum likelihood learning
 - $\max_{\psi} \frac{1}{N} \sum_{i=1}^N \log p(\tau_i|O_{1:T}, \psi) = \max_{\psi} \frac{1}{N} \sum_{i=1}^N r_\psi(\tau_i) - \log Z$
 - Where $Z = \int p(\tau) \exp(r_\psi(\tau)) d\tau$ is the partition function

Inverse Reinforcement Learning

The IRL partition function

► The IRL partition function

- $\max_{\psi} \frac{1}{N} \sum_{i=1}^N \log p(\tau_i | O_{1:T}, \psi) = \max_{\psi} \frac{1}{N} \sum_{i=1}^N r_{\psi}(\tau_i) - \log Z$
- $Z = \int p(\tau) \exp(r_{\psi}(\tau)) d\tau$
- Gradient
 - $\nabla_{\psi} L = \frac{1}{N} \sum_{i=1}^N \nabla_{\psi} r_{\psi}(\tau_i) - \frac{1}{Z} \int p(\tau) \exp(r_{\psi}(\tau)) \nabla_{\psi} r_{\psi}(\tau) d\tau$
 - Remember $p(\tau | O_{1:T}, \psi) \sim \frac{1}{Z} p(\tau) \exp(\sum_t r(s_t, a_t)) = \frac{1}{Z} p(\tau) \exp(r_{\psi}(\tau))$
 - Then we have
 - $\nabla_{\psi} L = E_{\tau \sim \pi^*(\tau)} [\nabla_{\psi} r_{\psi}(\tau_i)] - E_{\tau \sim p(\tau | O_{1:T}, \psi)} [\nabla_{\psi} r_{\psi}(\tau)]$
 - $E_{\tau \sim \pi^*(\tau)} [\nabla_{\psi} r_{\psi}(\tau_i)]$
 - estimate from expert samples
 - $E_{\tau \sim p(\tau | O_{1:T}, \psi)} [\nabla_{\psi} r_{\psi}(\tau)]$
 - soft optimal policy under current reward

Inverse Reinforcement Learning

The IRL partition function

► Estimating the expectation

- $E_{\tau \sim p(\tau | O_{1:T}, \psi)} [\nabla_{\psi} r_{\psi}(\tau)] = E_{\tau \sim p(\tau | O_{1:T}, \psi)} [\nabla_{\psi} \sum_{t=1}^T r_{\psi}(s_t, a_t)]$
- $= \sum_{i=1}^T E_{(s_t, a_t) \sim p(s_t, a_t | O_{1:T}, \psi)} [\nabla_{\psi} r_{\psi}(s_t, a_t)]$
- Where
 - $p(s_t, a_t | O_{1:T}, \psi) = p(a_t | s_t, O_{1:T}, \psi)p(s_t | O_{1:T}, \psi)$
- Remember that
 - $p(a_t | s_t, O_{1:T}, \psi) = \frac{\beta(s_t, a_t)}{\beta(s_t)}$
 - $p(s_t | O_{1:T}, \psi) \propto \beta_t(s_t)\alpha_t(s_t)$
- Then we have
 - $p(s_t, a_t | O_{1:T}, \psi) = p(a_t | s_t, O_{1:T}, \psi)p(s_t | O_{1:T}, \psi)$
 - $\propto \beta(s_t, a_t)\alpha_t(s_t)$
- Let $\mu(s_t, a_t) \propto \beta(s_t, a_t)\alpha_t(s_t)$
 - $E_{\tau \sim p(\tau | O_{1:T}, \psi)} [\nabla_{\psi} r_{\psi}(\tau)] = \sum_{t=1}^T \int \int \mu(s_t, a_t) \nabla_{\psi} r_{\psi}(s_t, a_t) ds_t da_t$
 - $= \sum_{t=1}^T \vec{\mu}_t^T \nabla_{\psi} \vec{r}_{\psi}$
 - Where $\vec{\mu}_t^T$ is state-action visit probability for each (s_t, a_t)

Inverse Reinforcement Learning

The MaxEnt IRL algorithm

► The MaxEnt IRL algorithm

- Repeat
 - Given ψ , compute backward message $\beta(s_t, a_t)$
 - Given ψ , compute forward message $\alpha_t(s_t)$
 - Compute $\mu(s_t, a_t) \propto \beta(s_t, a_t)\alpha_t(s_t)$
 - Evaluate $\nabla_\psi L = \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\psi r_\psi(s_{i,t}, a_{i,t}) - \sum_{t=1}^T \int \int \mu(s_t, a_t) \nabla_\psi r_\psi(s_t, a_t) ds_t da_t$
 - $\psi \leftarrow \psi + \eta \nabla_\psi L$

► Why it is max entropy?

- For $r_\psi(s, a) = \psi^T f(s, a)$
- It optimizes $\max_\psi H(\pi^{r_\psi})$ such that $E_{\pi^{r_\psi}}[f(s, a)] = E_{\pi^*}[f(s, a)]$

Inverse Reinforcement Learning

Deep Inverse Reinforcement Learning

► What about larger RL problems?

- MaxEnt IRL: probabilistic framework for learning reward functions
- Computing gradient requires enumerating state-action visitations for all states and actions
 - Only really viable for small, discrete state and action spaces
 - Amounts to a dynamic programming algorithm (exact forward backward inference)
- For deep IRL, we want two things:
 - Large and continuous state and action spaces
 - Effective learning under unknown dynamics

Inverse Reinforcement Learning

Unknown Dynamics & Large State/action Spaces

- ▶ Assume we don't know the dynamics, but we can sample, like in standard RL

- $\nabla_\psi L = \mathbb{E}_{\tau \sim \pi^*(\tau)}[\nabla_\psi r_\psi(\tau_i)] - \mathbb{E}_{\tau \sim p(\tau|O_{1:T}, \psi)}[\nabla_\psi r_\psi(\tau)]$
 - Learn $p(a_t|s_t, O_{1:T}, \psi)$ using any max-ent RL algorithm
 - $J(\theta) = \sum_t \mathbb{E}_{\pi(s_t, a_t)}[r_\psi(s_t, a_t)] + \mathbb{E}_{\pi(s_t)}[H(\pi(a|s_t))]$
 - Then run this policy to sample $\{\tau_j\}$
 - Then
 - $\nabla_\psi L \approx \frac{1}{N} \sum_{i=1}^N \nabla_\psi r_\psi(\tau_i) - \frac{1}{M} \sum_{j=1}^M \nabla_\psi r_\psi(\tau_j)$
 - $\frac{1}{N} \sum_{i=1}^N \nabla_\psi r_\psi(\tau_i)$ sum over expert samples
 - $\frac{1}{M} \sum_{j=1}^M \nabla_\psi r_\psi(\tau_j)$ sum over policy samples

Inverse Reinforcement Learning

Unknown Dynamics & Large State/action Spaces

► More efficient sample-based updates

- $\nabla_{\psi} L \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\psi} r_{\psi}(\tau_i) - \frac{1}{M} \sum_{j=1}^M \nabla_{\psi} r_{\psi}(\tau_j)$
- Use lazy policy optimization

► Problems

- Wrong distribution, estimator is now biased

► Solution

- Use importance sampling

- $\nabla_{\psi} L \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\psi} r_{\psi}(\tau_i) - \frac{1}{\sum_j w_j} \sum_{j=1}^M w_j \nabla_{\psi} r_{\psi}(\tau_j)$
- Where $w_j = \frac{p(\tau_j | O_{1:T}, \psi)}{\pi(\tau_j)} = \frac{p(\tau_j) \exp(r_{\psi}(\tau_j))}{\pi(\tau_j)}$
- $= \frac{p(s_1) \prod_t p(s_{t+1} | s_t, a_t) \exp(r_{\psi}(s_t, a_t))}{p(s_1) \prod_t p(s_{t+1} | s_t, a_t) \pi(a_t | s_t)} = \frac{\exp(\sum_t r_{\psi}(s_t, a_t))}{\prod_t \pi(a_t | s_t)}$

Inverse Reinforcement Learning


Importance Sampling

► Importance sampling

- $w_j = \frac{\exp(\sum_t r_\psi(s_t, a_t))}{\prod_t \pi(a_t | s_t)}$
- Which sampling distribution $\pi(\tau)$ is best?
 - Optimal IS distribution $q(x)$ for $E_{p(x)}[f(x)]$ is $q(x) \propto |f(x)|p(x)$
 - Then, optimal π is $\pi(\tau) \propto \exp(r_\psi(\tau))$
 - Max-ent optimal policy for r_ψ
- Each policy update w.r.t r_ψ bring closer to the optimal distribution

Inverse Reinforcement Learning


Guided Cost Learning



Inverse Reinforcement Learning

Inverse Reinforcement Learning


- ▶ Looks like a adversarial game



Inverse Reinforcement Learning

Generative Adversarial Networks

► Generative Adversarial Networks



Inverse Reinforcement Learning

Inverse RL as a GAN

► GAN discriminator

- $L_{\text{Discriminator}}(\psi) = \mathbb{E}_{x \sim p^*}[-\log D_\psi(x)] + \mathbb{E}_{x \sim p_\theta}[-\log(1 - D_\psi(x))]$
- $\psi = \underset{\psi}{\operatorname{argmax}} \frac{1}{N} \sum_{x \sim p^*} \log D_\psi(x) + \frac{1}{M} \sum_{x \sim p_\theta} \log(1 - D_\psi(x))$
- The best discriminator $D(x) = p_\psi(\text{real image}|x)$ is
 - $D^*(x) = \frac{p^*(x)}{p_\theta(x) + p^*(x)}$

► IRL

- Optimal policy approaches $\pi_\theta(\tau) \propto p(\tau)\exp(r_\psi(\tau))$
- Then let $p^*(x) \propto p(\tau)\exp(r_\psi(\tau))$
 - $D_\psi(\tau) = \frac{p(\tau) \frac{1}{Z} \exp(r_\psi(\tau))}{p_\theta(\tau) + p(\tau) \frac{1}{Z} \exp(r_\psi(\tau))} = \frac{p(\tau) \frac{1}{Z} \exp(r_\psi(\tau))}{p(\tau) \prod_t \pi_\theta(a_t|s_t) + p(\tau) \frac{1}{Z} \exp(r_\psi(\tau))}$
 - $= \frac{\frac{1}{Z} \exp(r_\psi(\tau))}{\prod_t \pi_\theta(a_t|s_t) + \frac{1}{Z} \exp(r_\psi(\tau))}$
 - $\psi = \underset{\psi}{\operatorname{argmax}} \mathbb{E}_{\tau \sim p^*}[\log D_\psi(\tau)] + \mathbb{E}_{\tau \sim \pi_\theta}[\log(1 - D_\psi(\tau))]$

Inverse Reinforcement Learning

Inverse RL as a GAN

► IRL

- Discriminator


- $L_{\text{Discriminator}}(\psi) = \mathbb{E}_{\tau \sim p^*}[-\log D_\psi(\tau)] + \mathbb{E}_{\tau \sim \pi_\theta}[-\log(1 - D_\psi(\tau))]$
- $\psi = \underset{\psi}{\operatorname{argmax}} \mathbb{E}_{\tau \sim p^*}[\log D_\psi(\tau)] + \mathbb{E}_{\tau \sim \pi_\theta}[\log(1 - D_\psi(\tau))]$
- Can be proved that optimize Z w.r.t. same objective as ψ

- Generator

- $L_{\text{Generator}}(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [\log(1 - D_\psi(\tau)) - \log D_\psi(\tau)]$
- $= \mathbb{E}_{\tau \sim \pi_\theta} [\log p_\theta(\tau) + \log Z - r_\psi(\tau)]$

Inverse Reinforcement Learning

Inverse RL as a GAN



Inverse Reinforcement Learning

Inverse RL as a GAN

► Can we just use a regular discriminator?

- $\psi = \underset{\psi}{\operatorname{argmax}} \mathbb{E}_{\tau \sim p^*} [\log D_\psi(\tau)] + \mathbb{E}_{\tau \sim \pi_\theta} [\log(1 - D_\psi(\tau))]$
- $D_\psi(\tau)$ = standard binary neural net classifier
- $\nabla_\theta L \approx \frac{1}{M} \sum_{j=1}^M \nabla_\theta \log \pi_\theta(\tau_j) \log D_\psi(\tau_j)$
- Often simpler to set up optimization, fewer moving parts
- Discriminator knows nothing at convergence
- Generally cannot reoptimize the “reward”

Inverse Reinforcement Learning

Summary

- ▶ **IRL: infer unknown reward from expert demonstrations**
- ▶ **MaxEnt IRL: infer reward by learning under the control-as-inference framework**
- ▶ **MaxEnt IRL with dynamic programming: simple and efficient, but requires small state space and known dynamics**
- ▶ **Sampling-based MaxEnt IRL: generate samples to estimate the partition function**
 - Guided cost learning algorithm
 - Connection to generative adversarial networks
 - Generative adversarial imitation learning (not IRL per se, but similar)

Inverse Reinforcement Learning

Suggested Readings

- ▶ **Manuel Watter, Embed to Control: a locally linear latent dynamics models for control from raw images**
- ▶ **M Zhang, SOLAR: deep structured latent representation for model-based reinforcement learning**
- ▶ **Todorov. (2006). Linearly solvable Markov decision problems**
- ▶ **Todorov. (2008). General duality between optimal control and estimation**
- ▶ **Kappen. (2009). Optimal control as a graphical model inference problem**
- ▶ **Ziebart. (2010). Modeling interaction via the principle of maximal causal entropy**
- ▶ **Rawlik, Toussaint, Vijaykumar. (2013). On stochastic optimal control and reinforcement learning by approximate inference**
- ▶ **Haarnoja*, Tang*, Abbeel, L. (2017). Reinforcement learning with deep energy based models**

Inverse Reinforcement Learning

Suggested Readings

- ▶ Nachum, Norouzi, Xu, Schuurmans. (2017). Bridging the gap between value and policy based reinforcement learning
- ▶ Schulman, Abbeel, Chen. (2017). Equivalence between policy gradients and soft Q-learning
- ▶ Haarnoja, Zhou, Abbeel, L. (2018). Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor
- ▶ Christodoulou, Petros. "Soft Actor-Critic for Discrete Action Settings."
- ▶ Sallans & Hinton. Using Free Energies to Represent Q-values in a Multiagent Reinforcement Learning Task
- ▶ O'Donoghue et al. Combining Policy Gradient and Q-Learning
- ▶ Levine. Reinforcement Learning and Control as Probabilistic Inference: Tutorial and Review

Inverse Reinforcement Learning

Suggested Readings

- ▶ **Abbeel & Ng: Apprenticeship learning via inverse reinforcement learning**
- ▶ **Ratliff et al: Maximum margin planning**
- ▶ **Ziebart et al. : Maximum Entropy Inverse Reinforcement Learning**
- ▶ **Wulfmeier et al. Maximum Entropy Deep Inverse Reinforcement Learning**
- ▶ **Wulfmeier et al. Watch This: Scalable Cost-Function Learning for Path Planning in Urban Environments**
- ▶ **Finn et al. Guided Cost Learning**
- ▶ **Finn, Christiano, et al. A Connection Between Generative Adversarial Networks, Inverse Reinforcement Learning, and Energy-Based Models**
- ▶ **Fu et al. Learning Robust Rewards with Adversarial Inverse Reinforcement Learning**
- ▶ **Ho & Ermon. Generative adversarial imitation learning**